



Red Hat Enterprise Linux 8

Using SELinux

Basic and advanced configuration of Security-Enhanced Linux (SELinux)

Red Hat Enterprise Linux 8 Using SELinux

Basic and advanced configuration of Security-Enhanced Linux (SELinux)

Legal Notice

Copyright © 2020 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

This title assists users and administrators in learning the basics and principles upon which SELinux functions and describes practical tasks to set up and configure various services.

Table of Contents

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION	3
CHAPTER 1. GETTING STARTED WITH SELINUX	4
1.1. INTRODUCTION TO SELINUX	4
1.2. BENEFITS OF RUNNING SELINUX	5
1.3. SELINUX EXAMPLES	6
1.4. SELINUX ARCHITECTURE AND PACKAGES	6
1.5. SELINUX STATES AND MODES	7
CHAPTER 2. CHANGING SELINUX STATES AND MODES	9
2.1. PERMANENT CHANGES IN SELINUX STATES AND MODES	9
2.2. CHANGING TO PERMISSIVE MODE	9
2.3. CHANGING TO ENFORCING MODE	10
2.4. ENABLING SELINUX ON SYSTEMS THAT PREVIOUSLY HAD IT DISABLED	11
2.5. DISABLING SELINUX	12
2.6. CHANGING SELINUX MODES AT BOOT TIME	13
CHAPTER 3. MANAGING CONFINED AND UNCONFINED USERS	15
3.1. CONFINED AND UNCONFINED USERS	15
3.2. SELINUX USER CAPABILITIES	16
3.3. ADDING A NEW USER AUTOMATICALLY MAPPED TO THE SELINUX UNCONFINED_U USER	17
3.4. ADDING A NEW USER AS AN SELINUX-CONFINED USER	18
3.5. CONFIGURING THE SYSTEM TO CONFINED SELINUX USERS	19
3.5.1. Confining regular users	19
3.5.2. Confining administrator users	20
3.5.2.1. Confining an administrator by mapping to sysadm_u	20
3.5.2.2. Confining an administrator using sudo and the sysadm_r role	21
3.5.3. Additional resources	22
3.6. ADDITIONAL RESOURCES	23
CHAPTER 4. CONFIGURING SELINUX FOR APPLICATIONS AND SERVICES WITH NON-STANDARD CONFIGURATIONS	24
4.1. CUSTOMIZING THE SELINUX POLICY FOR THE APACHE HTTP SERVER IN A NON-STANDARD CONFIGURATION	24
4.2. ADJUSTING THE POLICY FOR SHARING NFS AND CIFS VOLUMES USING SELINUX BOOLEANS	26
4.3. ADDITIONAL RESOURCES	27
CHAPTER 5. TROUBLESHOOTING PROBLEMS RELATED TO SELINUX	28
5.1. IDENTIFYING SELINUX DENIALS	28
5.2. ANALYZING SELINUX DENIAL MESSAGES	29
5.3. FIXING ANALYZED SELINUX DENIALS	30
5.4. SELINUX DENIALS IN THE AUDIT LOG	33
5.5. RELATED INFORMATION	34
CHAPTER 6. CREATING SELINUX POLICIES FOR CONTAINERS	35
6.1. INTRODUCTION TO THE UDICA SELINUX POLICY GENERATOR	35
6.2. CREATING AND USING AN SELINUX POLICY FOR A CUSTOM CONTAINER	35
6.3. ADDITIONAL RESOURCES	38
CHAPTER 7. DEPLOYING THE SAME SELINUX CONFIGURATION ON MULTIPLE SYSTEMS	39
7.1. INTRODUCTION TO THE SELINUX SYSTEM ROLE	39
7.2. USING THE SELINUX SYSTEM ROLE TO APPLY SELINUX SETTINGS ON MULTIPLE SYSTEMS	40
7.3. TRANSFERRING SELINUX SETTINGS TO ANOTHER SYSTEM WITH SEMANAGE	41

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

We appreciate your input on our documentation. Please let us know how we could make it better. To do so:

- For simple comments on specific passages:
 1. Make sure you are viewing the documentation in the *Multi-page HTML* format. In addition, ensure you see the **Feedback** button in the upper right corner of the document.
 2. Use your mouse cursor to highlight the part of text that you want to comment on.
 3. Click the **Add Feedback** pop-up that appears below the highlighted text.
 4. Follow the displayed instructions.
- For submitting more complex feedback, create a Bugzilla ticket:
 1. Go to the [Bugzilla](#) website.
 2. As the Component, use **Documentation**.
 3. Fill in the **Description** field with your suggestion for improvement. Include a link to the relevant part(s) of documentation.
 4. Click **Submit Bug**.

CHAPTER 1. GETTING STARTED WITH SELINUX

1.1. INTRODUCTION TO SELINUX

Security Enhanced Linux (SELinux) provides an additional layer of system security. SELinux fundamentally answers the question: *May <subject> do <action> to <object>?*, for example: *May a web server access files in users' home directories?*

The standard access policy based on the user, group, and other permissions, known as Discretionary Access Control (DAC), does not enable system administrators to create comprehensive and fine-grained security policies, such as restricting specific applications to only viewing log files, while allowing other applications to append new data to the log files.

SELinux implements Mandatory Access Control (MAC). Every process and system resource has a special security label called an *SELinux context*. A SELinux context, sometimes referred to as an *SELinux label*, is an identifier which abstracts away the system-level details and focuses on the security properties of the entity. Not only does this provide a consistent way of referencing objects in the SELinux policy, but it also removes any ambiguity that can be found in other identification methods. For example, a file can have multiple valid path names on a system that makes use of bind mounts.

The SELinux policy uses these contexts in a series of rules which define how processes can interact with each other and the various system resources. By default, the policy does not allow any interaction unless a rule explicitly grants access.




NOTE

Remember that SELinux policy rules are checked after DAC rules. SELinux policy rules are not used if DAC rules deny access first, which means that no SELinux denial is logged if the traditional DAC rules prevent the access.

SELinux contexts have several fields: user, role, type, and security level. The SELinux type information is perhaps the most important when it comes to the SELinux policy, as the most common policy rule which defines the allowed interactions between processes and system resources uses SELinux types and not the full SELinux context. SELinux types end with **_t**. For example, the type name for the web server is **httpd_t**. The type context for files and directories normally found in **/var/www/html/** is **httpd_sys_content_t**. The type contexts for files and directories normally found in **/tmp** and **/var/tmp/** is **tmp_t**. The type context for web server ports is **http_port_t**.

There is a policy rule that permits Apache (the web server process running as **httpd_t**) to access files and directories with a context normally found in **/var/www/html/** and other web server directories (**httpd_sys_content_t**). There is no allow rule in the policy for files normally found in **/tmp** and **/var/tmp/**, so access is not permitted. With SELinux, even if Apache is compromised, and a malicious script gains access, it is still not able to access the **/tmp** directory.

Figure 1.1. An example how can SELinux help to run Apache and MariaDB in a secure way.



RHEL_467048_0218

As the previous scheme shows, SELinux allows the Apache process running as **httpd_t** to access the **/var/www/html/** directory and it denies the same process to access the **/data/mysql/** directory because there is no allow rule for the **httpd_t** and **mysqld_db_t** type contexts). On the other hand, the MariaDB process running as **mysqld_t** is able to access the **/data/mysql/** directory and SELinux also correctly denies the process with the **mysqld_t** type to access the **/var/www/html/** directory labeled as **httpd_sys_content_t**.

Additional resources

For more information, see the following documentation:

- The **selinux(8)** man page and man pages listed by the **apropos selinux** command.
- Man pages listed by the **man -k _selinux** command when the **selinux-policy-doc** package is installed.
- [The SELinux Coloring Book](#) helps you to better understand SELinux basic concepts.
- [SELinux Wiki FAQ](#)

1.2. BENEFITS OF RUNNING SELINUX

SELinux provides the following benefits:

- All processes and files are labeled. SELinux policy rules define how processes interact with files, as well as how processes interact with each other. Access is only allowed if an SELinux policy rule exists that specifically allows it.
- Fine-grained access control. Stepping beyond traditional UNIX permissions that are controlled at user discretion and based on Linux user and group IDs, SELinux access decisions are based on all available information, such as an SELinux user, role, type, and, optionally, a security level.
- SELinux policy is administratively-defined and enforced system-wide.
- Improved mitigation for privilege escalation attacks. Processes run in domains, and are therefore separated from each other. SELinux policy rules define how processes access files and other processes. If a process is compromised, the attacker only has access to the normal functions of that process, and to files the process has been configured to have access to. For

example, if the Apache HTTP Server is compromised, an attacker cannot use that process to read files in user home directories, unless a specific SELinux policy rule was added or configured to allow such access.

- SELinux can be used to enforce data confidentiality and integrity, as well as protecting processes from untrusted inputs.

However, SELinux is not:

- antivirus software,
- replacement for passwords, firewalls, and other security systems,
- all-in-one security solution.

SELinux is designed to enhance existing security solutions, not replace them. Even when running SELinux, it is important to continue to follow good security practices, such as keeping software up-to-date, using hard-to-guess passwords, and firewalls.

1.3. SELINUX EXAMPLES

The following examples demonstrate how SELinux increases security:

- The default action is deny. If an SELinux policy rule does not exist to allow access, such as for a process opening a file, access is denied.
- SELinux can confine Linux users. A number of confined SELinux users exist in the SELinux policy. Linux users can be mapped to confined SELinux users to take advantage of the security rules and mechanisms applied to them. For example, mapping a Linux user to the SELinux **user_u** user, results in a Linux user that is not able to run unless configured otherwise set user ID (setuid) applications, such as **sudo** and **su**, as well as preventing them from executing potentially malicious files and applications in their home directory.
- Increased process and data separation. The concept of SELinux *domains* allows defining which processes can access certain files and directories. For example, when running SELinux, unless otherwise configured, an attacker cannot compromise a Samba server, and then use that Samba server as an attack vector to read and write to files used by other processes, such as MariaDB databases.
- SELinux helps mitigate the damage made by configuration mistakes. Domain Name System (DNS) servers often replicate information between each other in what is known as a zone transfer. Attackers can use zone transfers to update DNS servers with false information. When running the Berkeley Internet Name Domain (BIND) as a DNS server in Red Hat Enterprise Linux, even if an administrator forgets to limit which servers can perform a zone transfer, the default SELinux policy prevents zone files^[1] from being updated using zone transfers, by the BIND **named** daemon itself, and by other processes.


1.4. SELINUX ARCHITECTURE AND PACKAGES

SELinux is a Linux Security Module (LSM) that is built into the Linux kernel. The SELinux subsystem in the kernel is driven by a security policy which is controlled by the administrator and loaded at boot. All security-relevant, kernel-level access operations on the system are intercepted by SELinux and examined in the context of the loaded security policy. If the loaded policy allows the operation, it continues. Otherwise, the operation is blocked and the process receives an error.

SELinux decisions, such as allowing or disallowing access, are cached. This cache is known as the Access Vector Cache (AVC). When using these cached decisions, SELinux policy rules need to be checked less, which increases performance. Remember that SELinux policy rules have no effect if DAC rules deny access first. Raw audit messages are logged to the `/var/log/audit/audit.log` and they start with the **type=AVC** string.

In Red Hat Enterprise Linux 8, system services are controlled by the **systemd** daemon; **systemd** starts and stops all services, and users and processes communicate with **systemd** using the **systemctl** utility. The **systemd** daemon can consult the SELinux policy and check the label of the calling process and the label of the unit file that the caller tries to manage, and then ask SELinux whether or not the caller is allowed the access. This approach strengthens access control to critical system capabilities, which include starting and stopping system services.

The **systemd** daemon also works as an SELinux Access Manager. It retrieves the label of the process running **systemctl** or the process that sent a **D-Bus** message to **systemd**. The daemon then looks up the label of the unit file that the process wanted to configure. Finally, **systemd** can retrieve information from the kernel if the SELinux policy allows the specific access between the process label and the unit file label. This means a compromised application that needs to interact with **systemd** for a specific service can now be confined by SELinux. Policy writers can also use these fine-grained controls to confine administrators.



IMPORTANT

To avoid incorrect SELinux labeling and subsequent problems, ensure that you start services using a **systemctl start** command.

Red Hat Enterprise Linux 8 provides the following packages for working with SELinux:

- policies: **selinux-policy-targeted**, **selinux-policy-mls**
- tools: **policycoreutils**, **policycoreutils-gui**, **libselinux-utils**, **policycoreutils-python-utils**, **setools-console**, **checkpolicy**

1.5. SELINUX STATES AND MODES

SELinux can run in one of three modes: enforcing, permissive, or disabled.

- Enforcing mode is the default, and recommended, mode of operation; in enforcing mode SELinux operates normally, enforcing the loaded security policy on the entire system.
- In permissive mode, the system acts as if SELinux is enforcing the loaded security policy, including labeling objects and emitting access denial entries in the logs, but it does not actually deny any operations. While not recommended for production systems, permissive mode can be helpful for SELinux policy development and debugging.
- Disabled mode is strongly discouraged; not only does the system avoid enforcing the SELinux policy, it also avoids labeling any persistent objects such as files, making it difficult to enable SELinux in the future.

Use the **setenforce** utility to change between enforcing and permissive mode. Changes made with **setenforce** do not persist across reboots. To change to enforcing mode, enter the **setenforce 1** command as the Linux root user. To change to permissive mode, enter the **setenforce 0** command. Use the **getenforce** utility to view the current SELinux mode:

```
# getenforce
Enforcing
```

```
# setenforce 0
# getenforce
Permissive
```

```
# setenforce 1
# getenforce
Enforcing
```

In Red Hat Enterprise Linux, you can set individual domains to permissive mode while the system runs in enforcing mode. For example, to make the *httpd_t* domain permissive:

```
# semanage permissive -a httpd_t
```

Note that permissive domains are a powerful tool that can compromise security of your system. Red Hat recommends to use permissive domains with caution, for example, when debugging a specific scenario.

[1] Text files that include information, such as host name to IP address mappings, that are used by DNS servers.

CHAPTER 2. CHANGING SELINUX STATES AND MODES

When enabled, SELinux can run in one of two modes: enforcing or permissive. The following sections show how to permanently change into these modes.

2.1. PERMANENT CHANGES IN SELINUX STATES AND MODES

As discussed in [SELinux states and modes](#), SELinux can be enabled or disabled. When enabled, SELinux has two modes: enforcing and permissive.

Use the **getenforce** or **sestatus** commands to check in which mode SELinux is running. The **getenforce** command returns **Enforcing**, **Permissive**, or **Disabled**.

The **sestatus** command returns the SELinux status and the SELinux policy being used:

```
$ sestatus
SELinux status:                enabled
SELinuxfs mount:              /sys/fs/selinux
SELinux root directory:       /etc/selinux
Loaded policy name:            targeted
Current mode:                  enforcing
Mode from config file:         enforcing
Policy MLS status:             enabled
Policy deny_unknown status:    allowed
Memory protection checking:    actual (secure)
Max kernel policy version:     31
```

NOTE

When systems run SELinux in permissive mode, users and processes can label various file-system objects incorrectly. File-system objects created while SELinux is disabled are not labeled at all. This behavior causes problems when changing to enforcing mode because SELinux relies on correct labels of file-system objects.

To prevent incorrectly labeled and unlabeled files from causing problems, file systems are automatically relabeled when changing from the disabled state to permissive or enforcing mode. In permissive mode, use the **fixfiles -F onboot** command as root to create the **/.autorelabel** file containing the **-F** option to ensure that files are relabeled upon next reboot.

2.2. CHANGING TO PERMISSIVE MODE

Use the following procedure to permanently change SELinux mode to permissive. When SELinux is running in permissive mode, SELinux policy is not enforced. The system remains operational and SELinux does not deny any operations but only logs AVC messages, which can be then used for troubleshooting, debugging, and SELinux policy improvements. Each AVC is logged only once in this case.

Prerequisites

- The **selinux-policy-targeted**, **libselinux-utils**, and **policycoreutils** packages are installed on your system.
- The **selinux=0** or **enforcing=0** kernel parameters are not used.

Procedure

1. Open the **/etc/selinux/config** file in a text editor of your choice, for example:

```
# vi /etc/selinux/config
```

2. Configure the **SELINUX=permissive** option:

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
#   enforcing - SELinux security policy is enforced.
#   permissive - SELinux prints warnings instead of enforcing.
#   disabled - No SELinux policy is loaded.
SELINUX=permissive
# SELINUXTYPE= can take one of these two values:
#   targeted - Targeted processes are protected,
#   mls - Multi Level Security protection.
SELINUXTYPE=targeted
```

3. Restart the system:

```
# reboot
```

Verification steps

1. After the system restarts, confirm that the **getenforce** command returns **Permissive**:

```
$ getenforce
Permissive
```

2.3. CHANGING TO ENFORCING MODE

Use the following procedure to switch SELinux to enforcing mode. When SELinux is running in enforcing mode, it enforces the SELinux policy and denies access based on SELinux policy rules. In RHEL, enforcing mode is enabled by default when the system was initially installed with SELinux.

Prerequisites

- The **selinux-policy-targeted**, **libselinux-utils**, and **policycoreutils** packages are installed on your system.
- The **selinux=0** or **enforcing=0** kernel parameters are not used.

Procedure

1. Open the **/etc/selinux/config** file in a text editor of your choice, for example:

```
# vi /etc/selinux/config
```

2. Configure the **SELINUX=enforcing** option:

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
```

```
# enforcing - SELinux security policy is enforced.
# permissive - SELinux prints warnings instead of enforcing.
# disabled - No SELinux policy is loaded.
SELINUX=enforcing
# SELINUXTYPE= can take one of these two values:
# targeted - Targeted processes are protected,
# mls - Multi Level Security protection.
SELINUXTYPE=targeted
```

3. Save the change, and restart the system:

```
# reboot
```

On the next boot, SELinux relabels all the files and directories within the system and adds SELinux context for files and directories that were created when SELinux was disabled.

Verification steps

1. After the system restarts, confirm that the **getenforce** command returns **Enforcing**:

```
$ getenforce
Enforcing
```

NOTE

After changing to enforcing mode, SELinux may deny some actions because of incorrect or missing SELinux policy rules. To view what actions SELinux denies, enter the following command as root:

```
# ausearch -m AVC,USER_AVC,SELINUX_ERR,USER_SELINUX_ERR -ts today
```

Alternatively, with the **setroubleshoot-server** package installed, enter:

```
# grep "SELinux is preventing" /var/log/messages
```

If SELinux is active and the Audit daemon (**auditd**) is not running on your system, then search for certain SELinux messages in the output of the **dmesg** command:

```
# dmesg | grep -i -e type=1300 -e type=1400
```

See [Troubleshooting problems related to SELinux](#) for more information.

2.4. ENABLING SELINUX ON SYSTEMS THAT PREVIOUSLY HAD IT DISABLED

When you enable SELinux on systems that previously had it disabled, to avoid problems, such as systems unable to boot or process failures, follow this procedure:

Procedure

1. Enable SELinux in permissive mode. For more information, see [Changing to permissive mode](#).

2. Restart your system:

```
# reboot
```

3. Check for SELinux denial messages. For more information, see [Identifying SELinux denials](#).
4. If there are no denials, switch to enforcing mode. For more information, see [Changing SELinux modes at boot time](#).

Verification steps

1. After the system restarts, confirm that the **getenforce** command returns **Enforcing**:


```
$ getenforce
Enforcing
```

Additional resources

- To run custom applications with SELinux in enforcing mode, choose one of the following scenarios:
 - Run your application in the **unconfined_service_t** domain.
 - Write a new policy for your application. See the [Writing Custom SELinux Policy](#) Knowledgebase article for more information.
- Temporary changes in modes are covered in [SELinux states and modes](#).

2.5. DISABLING SELINUX

Use the following procedure to permanently disable SELinux.



IMPORTANT

When SELinux is disabled, SELinux policy is not loaded at all; it is not enforced and AVC messages are not logged. Therefore, all [benefits of running SELinux](#) are lost.

Red Hat strongly recommends to use permissive mode instead of permanently disabling SELinux. See [Changing to permissive mode](#) for more information about permissive mode.



WARNING

Disabling SELinux using the **SELINUX=disabled** option in the **/etc/selinux/config** results in a process in which the kernel boots with SELinux enabled and switches to disabled mode later in the boot process. Because memory leaks and race conditions causing kernel panics can occur, prefer disabling SELinux by adding the **selinux=0** parameter to the kernel command line as described in [Changing SELinux modes at boot time](#) if your scenario really requires to completely disable SELinux.

Procedure

1. Open the `/etc/selinux/config` file in a text editor of your choice, for example:

```
# vi /etc/selinux/config
```

2. Configure the **SELINUX=disabled** option:

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
#   enforcing - SELinux security policy is enforced.
#   permissive - SELinux prints warnings instead of enforcing.
#   disabled - No SELinux policy is loaded.
SELINUX=disabled
# SELINUXTYPE= can take one of these two values:
#   targeted - Targeted processes are protected,
#   mls - Multi Level Security protection.
SELINUXTYPE=targeted
```

3. Save the change, and restart your system:

```
# reboot
```

Verification steps

1. After reboot, confirm that the **getenforce** command returns **Disabled**:

```
$ getenforce
Disabled
```

2.6. CHANGING SELINUX MODES AT BOOT TIME

On boot, you can set several kernel parameters to change the way SELinux runs:

enforcing=0

Setting this parameter causes the machine to boot in permissive mode, which is useful when troubleshooting issues. Using permissive mode might be the only option to detect a problem if your file system is too corrupted. Moreover, in permissive mode the system continues to create the labels correctly. The AVC messages that are created in this mode can be different than in enforcing mode. In permissive mode, only the first denial is reported. However, in enforcing mode you might get a denial on reading a directory and an application stops. In permissive mode, you get the same AVC message, but the application continues reading files in the directory and you get an AVC for each denial in addition.

selinux=0

This parameter causes the kernel to not load any part of the SELinux infrastructure. The init scripts notice that the system booted with the **selinux=0** parameter and touch the `/.autorelabel` file. This causes the system to automatically relabel the next time you boot with SELinux enabled.



IMPORTANT

Red Hat does not recommend using the **selinux=0** parameter. To debug your system, prefer using permissive mode.

autorelabel=1

This parameter forces the system to relabel similarly to the following commands:

```
# touch /.autorelabel  
# reboot
```

If a file system contains a large amount of mislabeled objects, start the system in permissive mode to make the autorelabel process successful.

Additional resources

- For additional SELinux-related kernel boot parameters, such as **checkreqprot**, see the **/usr/share/doc/kernel-doc-*<KERNEL_VER>*/Documentation/admin-guide/kernel-parameters.txt** file installed with the **kernel-doc** package. Replace the *<KERNEL_VER>* string with the version number of the installed kernel, for example:

```
# yum install kernel-doc  
$ less /usr/share/doc/kernel-doc-4.18.0/Documentation/admin-guide/kernel-parameters.txt
```

CHAPTER 3. MANAGING CONFINED AND UNCONFINED USERS

The following sections explain the mapping of Linux users to SELinux users, describe the basic confined user domains, and demonstrate mapping a new user to an SELinux user.

3.1. CONFINED AND UNCONFINED USERS

Each Linux user is mapped to an SELinux user using SELinux policy. This allows Linux users to inherit the restrictions on SELinux users.

To see the SELinux user mapping on your system, use the **semanage login -l** command as root:

```
# semanage login -l
Login Name      SELinux User    MLS/MCS Range  Service
__default__     unconfined_u    s0-s0:c0.c1023 *
root            unconfined_u    s0-s0:c0.c1023 *
```

In Red Hat Enterprise Linux, Linux users are mapped to the SELinux **default** login by default, which is mapped to the SELinux **unconfined_u** user. The following line defines the default mapping:

```
__default__     unconfined_u    s0-s0:c0.c1023 *
```

Confined and unconfined Linux users are subject to executable and writable memory checks, and are also restricted by MCS or MLS.

To list the available SELinux users, enter the following command:

```
$ seinfo -u
Users: 8
  guest_u
  root
  staff_u
  sysadm_u
  system_u
  unconfined_u
  user_u
  xguest_u
```

Note that the **seinfo** command is provided by the **setools-console** package, which is not installed by default.

If an unconfined Linux user executes an application that SELinux policy defines as one that can transition from the **unconfined_t** domain to its own confined domain, the unconfined Linux user is still subject to the restrictions of that confined domain. The security benefit of this is that, even though a Linux user is running unconfined, the application remains confined. Therefore, the exploitation of a flaw in the application can be limited by the policy.

Similarly, we can apply these checks to confined users. Each confined user is restricted by a confined user domain. The SELinux policy can also define a transition from a confined user domain to its own target confined domain. In such a case, confined users are subject to the restrictions of that target confined domain. The main point is that special privileges are associated with the confined users according to their role.

3.2. SELINUX USER CAPABILITIES

The following table provides examples of basic confined domains for Linux users in Red Hat Enterprise Linux:

Table 3.1. SELinux user capabilities

User	Role	Domain	X Window System	su or sudo	Execute in home directory and /tmp (default)	Networking
sysadm_u	sysadm_r	sysadm_t	yes	su and sudo	yes	yes
staff_u	staff_r	staff_t	yes	only sudo	yes	yes
user_u	user_r	user_t	yes	no	yes	yes
guest_u	guest_r	guest_t	no	no	yes	no
xguest_u	xguest_r	xguest_t	yes	no	yes	Firefox only

- Linux users in the **user_t**, **guest_t**, and **xguest_t** domains can only run set user ID (setuid) applications if SELinux policy permits it (for example, **passwd**). These users cannot run the **su** and **sudo** setuid applications, and therefore cannot use these applications to become root.
- Linux users in the **sysadm_t**, **staff_t**, **user_t**, and **xguest_t** domains can log in using the X Window System and a terminal.
- By default, Linux users in the **staff_t**, **user_t**, **guest_t**, and **xguest_t** domains can execute applications in their home directories and **/tmp**. To prevent them from executing applications, which inherit users' permissions, in directories they have write access to, set the **guest_exec_content** and **xguest_exec_content** booleans to **off**. This helps prevent flawed or malicious applications from modifying users' files.
- The only network access Linux users in the **xguest_t** domain have is Firefox connecting to web pages.
- The **sysadm_u** user cannot log in directly using SSH. To enable SSH logins for **sysadm_u**, set the **ssh_sysadm_login** boolean to **on**:

```
# setsebool -P ssh_sysadm_login on
```

Note that **system_u** is a special user identity for system processes and objects. It must never be associated to a Linux user. Also, **unconfined_u** and **root** are unconfined users. For these reasons, they are not included in the previous table of SELinux user capabilities.

Alongside with the already mentioned SELinux users, there are special roles, that can be mapped to those users using the **semanage user** command. These roles determine what SELinux allows the user to do:

- webadm_r** can only administrate SELinux types related to the Apache HTTP Server.

- **dbadm_r** can only administrate SELinux types related to the MariaDB database and the PostgreSQL database management system.
- **logadm_r** can only administrate SELinux types related to the **syslog** and **auditlog** processes.
- **secadm_r** can only administrate SELinux.
- **auditadm_r** can only administrate processes related to the Audit subsystem.

To list all available roles, enter the the **seinfo -r** command:

```
$ seinfo -r
Roles: 14
  auditadm_r
  dbadm_r
  guest_r
  logadm_r
  nx_server_r
  object_r
  secadm_r
  staff_r
  sysadm_r
  system_r
  unconfined_r
  user_r
  webadm_r
  xguest_r
```

Note that the **seinfo** command is provided by the **setools-console** package, which is not installed by default.

Additional resources

- For more information, see the **seinfo(1)**, **semanage-login(8)**, and **xguest_selinux(8)** man pages.

3.3. ADDING A NEW USER AUTOMATICALLY MAPPED TO THE SELINUX UNCONFINED_U USER

The following procedure demonstrates how to add a new Linux user to the system. The user is automatically mapped to the SELinux **unconfined_u** user.

Prerequisites

- The **root** user is running unconfined, as it does by default in Red Hat Enterprise Linux.

Procedure

1. Enter the following command to create a new Linux user named *example.user*:

```
# useradd example.user
```

2. To assign a password to the Linux *example.user* user:

```
# passwd example.user
Changing password for user example.user.
New password:
Retype new password:
passwd: all authentication tokens updated successfully.
```

3. Log out of your current session.
4. Log in as the Linux *example.user* user. When you log in, the **pam_selinux** PAM module automatically maps the Linux user to an SELinux user (in this case, **unconfined_u**), and sets up the resulting SELinux context. The Linux user's shell is then launched with this context.

Verification steps

1. When logged in as the *example.user* user, check the context of a Linux user:

```
$ id -Z
unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023
```

Additional resources

- For more information, see the **pam_selinux(8)** man page.

3.4. ADDING A NEW USER AS AN SELINUX-CONFINED USER

Use the following steps to add a new SELinux-confined user to the system. This example procedure maps the user to the SELinux **staff_u** user right with the command for creating the user account.

Prerequisites

- The **root** user is running unconfined, as it does by default in Red Hat Enterprise Linux.

Procedure

1. Enter the following command to create a new Linux user named *example.user* and map it to the SELinux **staff_u** user:

```
# useradd -Z staff_u example.user
```

2. To assign a password to the Linux *example.user* user:

```
# passwd example.user
Changing password for user example.user.
New password:
Retype new password:
passwd: all authentication tokens updated successfully.
```

3. Log out of your current session.
4. Log in as the Linux *example.user* user. The user's shell launches with the **staff_u** context.

Verification steps

1. When logged in as the *example.user* user, check the context of a Linux user:

```
$ id -Z
uid=1000(example.user) gid=1000(example.user) groups=1000(example.user)
context=staff_u:staff_r:staff_t:s0-s0:c0.c1023
```

Additional resources

- For more information, see the **pam_selinux(8)** man page.

3.5. CONFIGURING THE SYSTEM TO CONFINED SELINUX USERS

By default, all Linux users in Red Hat Enterprise Linux, including users with administrative privileges, are mapped to the unconfined SELinux user **unconfined_u**. You can improve the security of the system by assigning users to SELinux confined users. This is useful to conform with the [V-71971 Security Technical Implementation Guide](#). For more information about confined and unconfined users, see [Managing confined and unconfined users](#).

3.5.1. Confining regular users

You can confine all regular users on your system by mapping them to the **user_u** SELinux user.

Procedure

1. Display the list of SELinux login records. The list displays the mappings of Linux users to SELinux users:

```
# semanage login -l

Login Name  SELinux User  MLS/MCS Range  Service
__default__ unconfined_u  s0-s0:c0.c1023  *
root        unconfined_u  s0-s0:c0.c1023  *
```

2. Map the **__default__** user, which represents all users without an explicit mapping, to the **user_u** SELinux user:

```
# semanage login -m -s user_u -r s0 __default__
```

Verification steps

1. Check that the **__default__** user is mapped to the **user_u** SELinux user:

```
# semanage login -l

Login Name  SELinux User  MLS/MCS Range  Service
__default__ user_u        s0              *
root        unconfined_u  s0-s0:c0.c1023  *
```

2. Verify that the processes of a new user run in the **user_u:user_r:user_t:s0** SELinux context.
 - a. Create a new user:

```
# adduser example.user
```

- b. Define a password for *example.user*:

```
# passwd example.user
```

- c. Log out as **root** and log in as the new user.

- d. Show the security context for the user's ID:

```
[example.user@localhost ~]$ id -Z
user_u:user_r:user_t:s0
```

- e. Show the security context of the user's current processes:

```
[example.user@localhost ~]$ ps axZ
LABEL                PID TTY   STAT  TIME COMMAND
-                    1 ?      Ss    0:05 /usr/lib/systemd/systemd --switched-root --
system --deserialize 18
-                   3729 ?      S     0:00 (sd-pam)
user_u:user_r:user_t:s0 3907 ?      Ss    0:00 /usr/lib/systemd/systemd --user
-                   3911 ?      S     0:00 (sd-pam)
user_u:user_r:user_t:s0 3918 ?      S     0:00 sshd: example.user@pts/0
user_u:user_r:user_t:s0 3922 pts/0    Ss    0:00 -bash
user_u:user_r:user_dbusd_t:s0 3969 ?      Ssl   0:00 /usr/bin/dbus-daemon --session --
address=systemd: --nofork --nopidfile --systemd-activation --syslog-only
user_u:user_r:user_t:s0 3971 pts/0    R+    0:00 ps axZ
```

3.5.2. Confining administrator users

You can use one of the following two methods to confine administrator users.

3.5.2.1. Confining an administrator by mapping to **sysadm_u**

You can confine a user with administrative privileges by mapping the user directly to the **sysadm_u** SELinux user. When the user logs in, the session runs in the **sysadm_u:sysadm_r:sysadm_t** SELinux context.

Prerequisites

- The **root** user runs unconfined. This is the Red Hat Enterprise Linux default.

Procedure

- Optional: To allow **sysadm_u** users to connect to the system using SSH:

```
# setsebool -P ssh_sysadm_login on
```

- Create a new user, add the user to the **wheel** user group, and map the user to the **sysadm_u** SELinux user:

```
# adduser -G wheel -Z sysadm_u example.user
```


- Optional: Map an existing user to the **sysadm_u** SELinux user and add the user to the **wheel** user group:

```
# usermod -G wheel -Z sysadm_u example.user
```

Verification steps

- Check that **example.user** is mapped to the **sysadm_u** SELinux user:

```
# semanage login -l | grep example.user
example.user sysadm_u s0-s0:c0.c1023 *
```

- Log in as **example.user**, for example, using SSH, and show the user's security context:

```
[example.user@localhost ~]$ id -Z
sysadm_u:sysadm_r:sysadm_t:s0-s0:c0.c1023
```

- Switch to the **root** user:

```
$ sudo -i
[sudo] password for example.user:
```

- Verify that the security context remains unchanged:

```
# id -Z
sysadm_u:sysadm_r:sysadm_t:s0-s0:c0.c1023
```

- Try an administrative task, for example, restarting the **sshd** service:

```
# systemctl restart sshd
```

If there is no output, the command finished successfully.

If the command does not finish successfully, it prints the following message:

```
Failed to restart sshd.service: Access denied
See system logs and 'systemctl status sshd.service' for details.
```

3.5.2.2. Confining an administrator using sudo and the sysadm_r role

You can map a specific user with administrative privileges to the **staff_u** SELinux user, and configure **sudo** so that the user can gain the **sysadm_r** SELinux administrator role. This role allows the user to perform administrative tasks without SELinux denials. When the user logs in, the session runs in the **staff_u:staff_r:staff_t** SELinux context, but when the user enters a command using **sudo**, the session changes to the **staff_u:sysadm_r:sysadm_t** context.

Prerequisites

- The **root** user runs unconfined. This is the Red Hat Enterprise Linux default.

Procedure

1. Create a new user, add the user to the **wheel** user group, and map the user to the **staff_u** SELinux user:

```
# adduser -G wheel -Z staff_u example.user
```

2. Optional: Map an existing user to the **staff_u** SELinux user and add the user to the **wheel** user group:

```
# usermod -G wheel -Z staff_u example.user
```

3. To allow *example.user* to gain the SELinux administrator role, create a new file in the **/etc/sudoers.d/** directory, for example:

```
# visudo -f /etc/sudoers.d/example.user
```

4. Add the following line to the new file:

```
example.user ALL=(ALL) TYPE=sysadm_t ROLE=sysadm_r ALL
```

Verification steps

1. Check that **example.user** is mapped to the **staff_u** SELinux user:

```
# semanage login -l | grep example.user
example.user  staff_u  s0-s0:c0.c1023  *
```

2. Log in as *example.user*, for example, using SSH, and switch to the **root** user:

```
[example.user@localhost ~]$ sudo -i
[sudo] password for example.user:
```

3. Show the **root** security context:

```
# id -Z
staff_u:sysadm_r:sysadm_t:s0-s0:c0.c1023
```

4. Try an administrative task, for example, restarting the **sshd** service:

```
# systemctl restart sshd
```

If there is no output, the command finished successfully.

If the command does not finish successfully, it prints the following message:

```
Failed to restart sshd.service: Access denied
See system logs and 'systemctl status sshd.service' for details.
```

3.5.3. Additional resources

- For additional options, see the [How to set up a system with SELinux confined users](#) knowledgebase article.

- For more information, see the **user_selinux(8)**, **staff_selinux(8)**, and **sysadm_selinux(8)** man pages.

3.6. ADDITIONAL RESOURCES

- For more information, see the **unconfined_selinux(8)** man page.

CHAPTER 4. CONFIGURING SELINUX FOR APPLICATIONS AND SERVICES WITH NON-STANDARD CONFIGURATIONS

When SELinux is in enforcing mode, the default policy is the targeted policy. The following sections provide information on setting up and configuring the SELinux policy for various services after you change configuration defaults, such as ports, database locations, or file-system permissions for processes.

In the following procedures, you learn to change SELinux types for non-standard ports, to identify and fix incorrect labels for changes of default directories, and to adjust the policy using SELinux booleans.

4.1. CUSTOMIZING THE SELINUX POLICY FOR THE APACHE HTTP SERVER IN A NON-STANDARD CONFIGURATION

You can configure the Apache HTTP server to listen on a different port and to provide content in a non-default directory. To prevent consequent SELinux denials, follow the steps in this procedure to adjust your system's SELinux policy.

Prerequisites

- The **httpd** package is installed and the Apache HTTP server is configured to listen on TCP port 3131 and to use the **/var/test_www/** directory instead of the default **/var/www/** directory.
- The **polycoreutils-python-utils** and **setroubleshoot-server** packages are installed on your system.

Procedure

1. Start the **httpd** service and check the status:

```
# systemctl start httpd
# systemctl status httpd
...
httpd[14523]: (13)Permission denied: AH00072: make_sock: could not bind to address
[::]:3131
...
systemd[1]: Failed to start The Apache HTTP Server.
...
```

2. The SELinux policy assumes that **httpd** runs on port 80:

```
# semanage port -l | grep http
http_cache_port_t      tcp      8080, 8118, 8123, 10001-10010
http_cache_port_t      udp      3130
http_port_t            tcp      80, 81, 443, 488, 8008, 8009, 8443, 9000
pegasus_http_port_t    tcp      5988
pegasus_https_port_t   tcp      5989
```

3. Change the SELinux type of port 3131 to match port 80:

```
# semanage port -a -t http_port_t -p tcp 3131
```

4. Start **httpd** again:

```
# systemctl start httpd
```

- However, the content remains inaccessible:

```
# wget localhost:3131/index.html
...
HTTP request sent, awaiting response... 403 Forbidden
...
```

Find the reason with the **sealert** tool:

```
# sealert -l ""
...
SELinux is preventing httpd from getattr access on the file /var/test_www/html/index.html.
...
```

- Compare SELinux types for the standard and the new path using the **matchpathcon** tool:

```
# matchpathcon /var/www/html /var/www/html system_u:object_r:httpd_sys_content_t:s0
# matchpathcon /var/test_www/html /var/test_www/html system_u:object_r:var_t:s0
```

- Change the SELinux type of the new **/var/test_www/html/** content directory to the type of the default **/var/www/html** directory:

```
# semanage fcontext -a -e /var/www /var/test_www
```

- Relabel the **/var** directory recursively:

```
# restorecon -Rv /var/
...
Relabeled /var/test_www/html from unconfined_u:object_r:var_t:s0 to
unconfined_u:object_r:httpd_sys_content_t:s0
Relabeled /var/test_www/html/index.html from unconfined_u:object_r:var_t:s0 to
unconfined_u:object_r:httpd_sys_content_t:s0
```

Verification steps

- Check that the **httpd** service is running:

```
# systemctl status httpd
...
Active: active (running)
...
systemd[1]: Started The Apache HTTP Server.
httpd[14888]: Server configured, listening on: port 3131
...
```

- Verify that the content provided by the Apache HTTP server is accessible:

```
# wget localhost:3131/index.html
...
HTTP request sent, awaiting response... 200 OK
```

```
Length: 0 [text/html]
Saving to: 'index.html'
...
```

Additional resources

- The **semanage(8)**, **matchpathcon(8)**, and **sealert(8)** man pages.

4.2. ADJUSTING THE POLICY FOR SHARING NFS AND CIFS VOLUMES USING SELINUX BOOLEANS

You can change parts of SELinux policy at runtime using booleans, even without any knowledge of SELinux policy writing. This enables changes, such as allowing services access to NFS volumes, without reloading or recompiling SELinux policy. The following procedure demonstrates listing SELinux booleans and configuring them to achieve the required changes in the policy.

NFS mounts on the client side are labeled with a default context defined by a policy for NFS volumes. In RHEL, this default context uses the **nfs_t** type. Also, Samba shares mounted on the client side are labeled with a default context defined by the policy. This default context uses the **cifs_t** type. You can enable or disable booleans to control which services are allowed to access the **nfs_t** and **cifs_t** types.

To allow the Apache HTTP server service (**httpd**) to access and share NFS and CIFS volumes, perform the following steps:

Prerequisites

- Optionally, install the **selinux-policy-devel** package to obtain clearer and more detailed descriptions of SELinux booleans in the output of the **semanage boolean -l** command.

Procedure

1. Identify SELinux booleans relevant for NFS, CIFS, and Apache:

```
# semanage boolean -l | grep 'nfs\|cifs' | grep httpd
httpd_use_cifs      (off , off) Allow httpd to access cifs file systems
httpd_use_nfs       (off , off) Allow httpd to access nfs file systems
```

2. List the current state of the booleans:

```
$ getsebool -a | grep 'nfs\|cifs' | grep httpd
httpd_use_cifs --> off
httpd_use_nfs  --> off
```

3. Enable the identified booleans:

```
# setsebool httpd_use_nfs on
# setsebool httpd_use_cifs on
```



NOTE

Use **setsebool** with the **-P** option to make the changes persistent across restarts. A **setsebool -P** command requires a rebuild of the entire policy, and it might take some time depending on your configuration.

Verification steps

1. Check that the booleans are **on**:

```
$ getsebool -a | grep 'nfs\|cifs' | grep httpd
httpd_use_cifs --> on
httpd_use_nfs --> on
```

Additional resources

- The **semanage-boolean(8)**, **sepolicy-booleans(8)**, **getsebool(8)**, **setsebool(8)**, **booleans(5)**, and **booleans(8)** man pages.

4.3. ADDITIONAL RESOURCES

- See [Troubleshooting problems related to SELinux](#) for more details on identifying and analyzing SELinux denials.

CHAPTER 5. TROUBLESHOOTING PROBLEMS RELATED TO SELINUX

If you plan to enable SELinux on systems where it has been previously disabled or if you run a service in a non-standard configuration, you might need to troubleshoot situations potentially blocked by SELinux. Note that in most cases, SELinux denials are signs of misconfiguration.

5.1. IDENTIFYING SELINUX DENIALS

Follow only the necessary steps from this procedure; in most cases, you need to perform just step 1.

Procedure

1. When your scenario is blocked by SELinux, the `/var/log/audit/audit.log` file is the first place to check for more information about a denial. To query Audit logs, use the **ausearch** tool. Because the SELinux decisions, such as allowing or disallowing access, are cached and this cache is known as the Access Vector Cache (AVC), use the **AVC** and **USER_AVC** values for the message type parameter, for example:

```
# ausearch -m AVC,USER_AVC,SELINUX_ERR,USER_SELINUX_ERR -ts recent
```

If there are no matches, check if the Audit daemon is running. If it does not, repeat the denied scenario after you start **auditd** and check the Audit log again.

2. In case **auditd** is running, but there are no matches in the output of **ausearch**, check messages provided by the **systemd** Journal:

```
# journalctl -t setroubleshoot
```

3. If SELinux is active and the Audit daemon is not running on your system, then search for certain SELinux messages in the output of the **dmesg** command:

```
# dmesg | grep -i -e type=1300 -e type=1400
```

4. Even after the previous three checks, it is still possible that you have not found anything. In this case, AVC denials can be silenced because of **dontaudit** rules.

To temporarily disable **dontaudit** rules, allowing all denials to be logged:

```
# semodule -DB
```

After re-running your denied scenario and finding denial messages using the previous steps, the following command enables **dontaudit** rules in the policy again:

```
# semodule -B
```

5. If you apply all four previous steps, and the problem still remains unidentified, consider if SELinux really blocks your scenario:

- Switch to permissive mode:


```
# setenforce 0
$ getenforce
Permissive
```

- Repeat your scenario.

If the problem still occurs, something different than SELinux is blocking your scenario.

5.2. ANALYZING SELINUX DENIAL MESSAGES

After [identifying](#) that SELinux is blocking your scenario, you might need to analyze the root cause before you choose a fix.

Prerequisites

- The **policycoreutils-python-utils** and **setroubleshoot-server** packages are installed on your system.

Procedure

1. List more details about a logged denial using the **sealert** command, for example:

```
$ sealert -l ""
SELinux is preventing /usr/bin/passwd from write access on the file
/root/test.

***** Plugin leaks (86.2 confidence) suggests *****

If you want to ignore passwd trying to write access the test file,
because you believe it should not need this access.
Then you should report this as a bug.
You can generate a local policy module to dontaudit this access.
Do
# ausearch -x /usr/bin/passwd --raw | audit2allow -D -M my-passwd
# semodule -X 300 -i my-passwd.pp

***** Plugin catchall (14.7 confidence) suggests *****

...

Raw Audit Messages
type=AVC msg=audit(1553609555.619:127): avc: denied { write } for
pid=4097 comm="passwd" path="/root/test" dev="dm-0" ino=17142697
scontext=unconfined_u:unconfined_r:passwd_t:s0-s0:c0.c1023
tcontext=unconfined_u:object_r:admin_home_t:s0 tclass=file permissive=0

...

Hash: passwd,passwd_t,admin_home_t,file,write
```

2. If the output obtained in the previous step does not contain clear suggestions:

- Enable full-path auditing to see full paths to accessed objects and to make additional Linux Audit event fields visible:

—

```
# auditctl -w /etc/shadow -p w -k shadow-write
```

- Clear the **setroubleshoot** cache:

```
# rm -f /var/lib/setroubleshoot/setroubleshoot.xml
```

- Reproduce the problem.
 - Repeat step 1.
3. If **sealert** returns only **catchall** suggestions or suggests adding a new rule using the **audit2allow** tool, match your problem with examples listed and explained in [SELinux denials in the Audit log](#).

Additional resources

- The **sealert(8)** man page.

5.3. FIXING ANALYZED SELINUX DENIALS

In most cases, suggestions provided by the **sealert** tool give you the right guidance about how to fix problems related to the SELinux policy. See [Analyzing SELinux denial messages](#) for information how to use **sealert** to analyze SELinux denials.

Be careful when the tool suggests using the **audit2allow** tool for configuration changes. You should not use **audit2allow** to generate a local policy module as your first option when you see an SELinux denial. Troubleshooting should start with a check if there is a labeling problem. The second most often case is that you have changed a process configuration, and you forgot to tell SELinux about it.

Labeling problems

A common cause of labeling problems is when a non-standard directory is used for a service. For example, instead of using **/var/www/html/** for a website, an administrator might want to use **/srv/myweb/**. On Red Hat Enterprise Linux, the **/srv** directory is labeled with the **var_t** type. Files and directories created in **/srv** inherit this type. Also, newly-created objects in top-level directories, such as **/myserver**, can be labeled with the **default_t** type. SELinux prevents the Apache HTTP Server (**httpd**) from accessing both of these types. To allow access, SELinux must know that the files in **/srv/myweb/** are to be accessible by **httpd**:

```
# semanage fcontext -a -t httpd_sys_content_t "/srv/myweb(/.*)?"
```

This **semanage** command adds the context for the **/srv/myweb/** directory and all files and directories under it to the SELinux file-context configuration. The **semanage** utility does not change the context. As root, use the **restorecon** utility to apply the changes:

```
# restorecon -R -v /srv/myweb
```

Incorrect context

The **matchpathcon** utility checks the context of a file path and compares it to the default label for that path. The following example demonstrates the use of **matchpathcon** on a directory that contains incorrectly labeled files:

```
$ matchpathcon -V /var/www/html/*
/var/www/html/index.html has context unconfined_u:object_r:user_home_t:s0, should be
```

```
system_u:object_r:httpd_sys_content_t:s0
/var/www/html/page1.html has context unconfined_u:object_r:user_home_t:s0, should be
system_u:object_r:httpd_sys_content_t:s0
```

In this example, the **index.html** and **page1.html** files are labeled with the **user_home_t** type. This type is used for files in user home directories. Using the **mv** command to move files from your home directory may result in files being labeled with the **user_home_t** type. This type should not exist outside of home directories. Use the **restorecon** utility to restore such files to their correct type:

```
# restorecon -v /var/www/html/index.html
restorecon reset /var/www/html/index.html context unconfined_u:object_r:user_home_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

To restore the context for all files under a directory, use the **-R** option:

```
# restorecon -R -v /var/www/html/
restorecon reset /var/www/html/page1.html context unconfined_u:object_r:samba_share_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
restorecon reset /var/www/html/index.html context unconfined_u:object_r:samba_share_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

Confined applications configured in non-standard ways

Services can be run in a variety of ways. To account for that, you need to specify how you run your services. You can achieve this through SELinux booleans that allow parts of SELinux policy to be changed at runtime. This enables changes, such as allowing services access to NFS volumes, without reloading or recompiling SELinux policy. Also, running services on non-default port numbers requires policy configuration to be updated using the **semanage** command.

For example, to allow the Apache HTTP Server to communicate with MariaDB, enable the **httpd_can_network_connect_db** boolean:

```
# setsebool -P httpd_can_network_connect_db on
```

Note that the **-P** option makes the setting persistent across reboots of the system.

If access is denied for a particular service, use the **getsebool** and **grep** utilities to see if any booleans are available to allow access. For example, use the **getsebool -a | grep ftp** command to search for FTP related booleans:

```
$ getsebool -a | grep ftp
ftpd_anon_write --> off
ftpd_full_access --> off
ftpd_use_cifs --> off
ftpd_use_nfs --> off

ftpd_connect_db --> off
httpd_enable_ftp_server --> off
tftp_anon_write --> off
```

To get a list of booleans and to find out if they are enabled or disabled, use the **getsebool -a** command. To get a list of booleans including their meaning, and to find out if they are enabled or disabled, install the **selinux-policy-devel** package and use the **semanage boolean -l** command as root.

Port numbers

Depending on policy configuration, services can only be allowed to run on certain port numbers. Attempting to change the port a service runs on without changing policy may result in the service failing to start. For example, run the **semanage port -l | grep http** command as root to list **http** related ports:

```
# semanage port -l | grep http
http_cache_port_t      tcp    3128, 8080, 8118
http_cache_port_t      udp    3130
http_port_t            tcp    80, 443, 488, 8008, 8009, 8443
pegasus_http_port_t    tcp    5988
pegasus_https_port_t   tcp    5989
```

The **http_port_t** port type defines the ports Apache HTTP Server can listen on, which in this case, are TCP ports 80, 443, 488, 8008, 8009, and 8443. If an administrator configures **httpd.conf** so that **httpd** listens on port 9876 (**Listen 9876**), but policy is not updated to reflect this, the following command fails:

```
# systemctl start httpd.service
Job for httpd.service failed. See 'systemctl status httpd.service' and 'journalctl -xn' for details.

# systemctl status httpd.service
httpd.service - The Apache HTTP Server
   Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
   Active: failed (Result: exit-code) since Thu 2013-08-15 09:57:05 CEST; 59s ago
   Process: 16874 ExecStop=/usr/sbin/httpd $OPTIONS -k graceful-stop (code=exited, status=0/SUCCESS)
   Process: 16870 ExecStart=/usr/sbin/httpd $OPTIONS -DFOREGROUND (code=exited, status=1/FAILURE)
```

An SELinux denial message similar to the following is logged to **/var/log/audit/audit.log**:

```
type=AVC msg=audit(1225948455.061:294): avc: denied { name_bind } for pid=4997
comm="httpd" src=9876 scontext=unconfined_u:system_r:httpd_t:s0
tcontext=system_u:object_r:port_t:s0 tclass=tcp_socket
```

To allow **httpd** to listen on a port that is not listed for the **http_port_t** port type, use the **semanage port** command to assign a different label to the port:

```
# semanage port -a -t http_port_t -p tcp 9876
```

The **-a** option adds a new record; the **-t** option defines a type; and the **-p** option defines a protocol. The last argument is the port number to add.

Corner cases, evolving or broken applications, and compromised systems

Applications may contain bugs, causing SELinux to deny access. Also, SELinux rules are evolving – SELinux may not have seen an application running in a certain way, possibly causing it to deny access, even though the application is working as expected. For example, if a new version of PostgreSQL is released, it may perform actions the current policy does not account for, causing access to be denied, even though access should be allowed.

For these situations, after access is denied, use the **audit2allow** utility to create a custom policy module to allow access. You can report missing rules in the SELinux policy in [Red Hat Bugzilla](#). For Red Hat Enterprise Linux 8, create bugs against the **Red Hat Enterprise Linux 8** product, and select the **selinux-policy** component. Include the output of the **audit2allow -w -a** and **audit2allow -a** commands in such bug reports.

If an application asks for major security privileges, it could be a signal that the application is compromised. Use intrusion detection tools to inspect such suspicious behavior.

The [Solution Engine](#) on the [Red Hat Customer Portal](#) can also provide guidance in the form of an article containing a possible solution for the same or very similar problem you have. Select the relevant product and version and use SELinux-related keywords, such as *selinux* or *avc*, together with the name of your blocked service or application, for example: **selinux samba**.

5.4. SELINUX DENIALS IN THE AUDIT LOG

The Linux Audit system stores log entries in the `/var/log/audit/audit.log` file by default. To list only SELinux-related records, use the **ausearch** command with the message type parameter set to **AVC** and **AVC_USER** at a minimum, for example:

```
# ausearch -m AVC,USER_AVC,SELINUX_ERR,USER_SELINUX_ERR
```

An SELinux denial entry in the Audit log file can look as follows:

```
type=AVC msg=audit(1395177286.929:1638): avc: denied { read } for pid=6591 comm="httpd"
name="webpages" dev="0:37" ino=2112 scontext=system_u:system_r:httpd_t:s0
tcontext=system_u:object_r:nfs_t:s0 tclass=dir
```

The most important parts of this entry are:

- **avc: denied** - the action performed by SELinux and recorded in Access Vector Cache (AVC)
- **{ read }** - the denied action
- **pid=6591** - the process identifier of the subject that tried to perform the denied action
- **comm="httpd"** - the name of the command that was used to invoke the analyzed process
- **httpd_t** - the SELinux type of the process
- **nfs_t** - the SELinux type of the object affected by the process action
- **tclass=dir** - the target object class

The previous log entry can be translated to:

*SELinux denied the **httpd** process with PID 6591 and the **httpd_t** type to read from a directory with the **nfs_t** type.*

The following SELinux denial message occurs when the Apache HTTP Server attempts to access a directory labeled with a type for the Samba suite:

```
type=AVC msg=audit(1226874073.147:96): avc: denied { getattr } for pid=2465 comm="httpd"
path="/var/www/html/file1" dev=dm-0 ino=284133 scontext=unconfined_u:system_r:httpd_t:s0
tcontext=unconfined_u:object_r:samba_share_t:s0 tclass=file
```

- **{ getattr }** - the **getattr** entry indicates the source process was trying to read the target file's status information. This occurs before reading files. SELinux denies this action because the process accesses the file and it does not have an appropriate label. Commonly seen permissions include **getattr**, **read**, and **write**.

- **path="/var/www/html/file1"** - the path to the object (target) the process attempted to access.
- **scontext="unconfined_u:system_r:httpd_t:s0"** - the SELinux context of the process (source) that attempted the denied action. In this case, it is the SELinux context of the Apache HTTP Server, which is running with the **httpd_t** type.
- **tcontext="unconfined_u:object_r:samba_share_t:s0"** - the SELinux context of the object (target) the process attempted to access. In this case, it is the SELinux context of **file1**.

This SELinux denial can be translated to:

*SELinux denied the **httpd** process with PID 2465 to access the **/var/www/html/file1** file with the **samba_share_t** type, which is not accessible to processes running in the **httpd_t** domain unless configured otherwise.*

Additional resources

- For more information, see the **auditd(8)** and **ausearch(8)** man pages.

5.5. RELATED INFORMATION

- The [Basic SELinux Troubleshooting in CLI](#) article on the Customer Portal.
- The [What is SELinux trying to tell me? The 4 key causes of SELinux errors](#) presentation on Fedora People

CHAPTER 6. CREATING SELINUX POLICIES FOR CONTAINERS

RHEL 8 provides a tool for generating SELinux policies for containers using the **udica** package. With **udica**, you can create a tailored security policy for better control of how a container accesses host system resources, such as storage, devices, and network. This enables you to harden your container deployments against security violations and it also simplifies achieving and maintaining regulatory compliance.

6.1. INTRODUCTION TO THE UDICA SELINUX POLICY GENERATOR

To simplify creating new SELinux policies for custom containers, RHEL 8 provides the **udica** utility. You can use this tool to create a policy based on an inspection of the container JavaScript Object Notation (JSON) file, which contains Linux-capabilities, mount-points, and ports definitions. The tool consequently combines rules generated using the results of the inspection with rules inherited from a specified SELinux Common Intermediate Language (CIL) block.

The process of generating SELinux policy for a container using **udica** has three main parts:

1. Parsing the container spec file in the JSON format
2. Finding suitable allow rules based on the results of the first part
3. Generating final SELinux policy

During the parsing phase, **udica** looks for Linux capabilities, network ports, and mount points.

Based on the results, **udica** detects which Linux capabilities are required by the container and creates an SELinux rule allowing all these capabilities. If the container binds to a specific port, **udica** uses SELinux user-space libraries to get the correct SELinux label of a port that is used by the inspected container.

Afterward, **udica** detects which directories are mounted to the container file-system name space from the host.

The CIL's block inheritance feature allows **udica** to create templates of SELinux *allow rules* focusing on a specific action, for example:

- *allow accessing home directories*
- *allow accessing log files*
- *allow accessing communication with Xserver .*

These templates are called blocks and the final SELinux policy is created by merging the blocks.

Additional resources

- For more details on the process of generating an SELinux policy with **udica**, see the [Generate SELinux policies for containers with udica](#) Red Hat Blog article.

6.2. CREATING AND USING AN SELINUX POLICY FOR A CUSTOM CONTAINER

To generate an SELinux security policy for a custom container, follow the steps in this procedure.

Prerequisites

- The **podman** tool for managing containers is installed. If it is not, use the **yum install podman** command.
- A custom Linux container - *ubi8* in this example.

Procedure

1. Install the **udica** package:

```
# yum install -y udica
```

2. Start the *ubi8* container that mounts the **/home** directory with read-only permissions and the **/var/spool** directory with permissions to read and write. The container exposes the port **21**.

```
# podman run --env container=podman -v /home:/home:ro -v /var/spool:/var/spool:rw -p 21:21 -it ubi8 bash
```

Note that now the container runs with the **container_t** SELinux type. This type is a generic domain for all containers in the SELinux policy and it might be either too strict or too loose for your scenario.

3. Enter the **podman ps** command to obtain the ID of the container:

```
# podman ps
CONTAINER ID  IMAGE                                COMMAND  CREATED      STATUS
PORTS  NAMES
37a3635afb8f  registry.access.redhat.com/ubi8:latest  bash    15 minutes ago  Up 15
minutes ago    heuristic_lewin
```

4. Create a container JSON file, and use **udica** for creating a policy module based on the information in the JSON file:

```
# podman inspect 37a3635afb8f > container.json
# udica -j container.json my_container
Policy my_container with container id 37a3635afb8f created!
[...]
```

Alternatively:

```
# podman inspect 37a3635afb8f | udica my_container
Policy my_container with container id 37a3635afb8f created!

Please load these modules using:
# semodule -i my_container.cil
/usr/share/udica/templates/{base_container.cil,net_container.cil,home_container.cil}

Restart the container with: "--security-opt label=type:my_container.process" parameter
```

5. As suggested by the output of **udica** in the previous step, load the policy module:


```
# semodule -i my_container.cil
/usr/share/udica/templates/{base_container.cil,net_container.cil,home_container.cil}
```

6. Stop the container and start it again with the **--security-opt label=type:my_container.process** option:

```
# podman stop 37a3635afb8f
# podman run --security-opt label=type:my_container.process -v /home:/home:ro -v
/var/spool:/var/spool:rw -p 21:21 -it ubi8 bash
```

Verification steps

1. Check that the container runs with the **my_container.process** type:

```
# ps -efZ | grep my_container.process
unconfined_u:system_r:container_runtime_t:s0-s0:c0.c1023 root 2275 434 1 13:49 pts/1
00:00:00 podman run --security-opt label=type:my_container.process -v /home:/home:ro -v
/var/spool:/var/spool:rw -p 21:21 -it ubi8 bash
system_u:system_r:my_container.process:s0:c270,c963 root 2317 2305 0 13:49 pts/0
00:00:00 bash
```

2. Verify that SELinux now allows access the **/home** and **/var/spool** mount points:

```
[root@37a3635afb8f /]# cd /home
[root@37a3635afb8f home]# ls
username
[root@37a3635afb8f ~]# cd /var/spool/
[root@37a3635afb8f spool]# touch test
[root@37a3635afb8f spool]#
```

3. Check that SELinux allows binding only to the port 21:

```
[root@37a3635afb8f /]# yum install nmap-ncat
[root@37a3635afb8f /]# nc -lvp 21
Ncat: Version 7.60 ( https://nmap.org/ncat )
Ncat: Generating a temporary 1024-bit RSA key. Use --ssl-key and --ssl-cert to use a
permanent one.
Ncat: SHA-1 fingerprint: 6EEC 102E 6666 5F96 CC4F E5FA A1BE 4A5E 6C76 B6DC
Ncat: Listening on :::21
Ncat: Listening on 0.0.0.0:21

[root@37a3635afb8f /]# nc -lvp 80
Ncat: Version 7.60 ( https://nmap.org/ncat )
Ncat: Generating a temporary 1024-bit RSA key. Use --ssl-key and --ssl-cert to use a
permanent one.
Ncat: SHA-1 fingerprint: 6EEC 102E 6666 5F96 CC4F E5FA A1BE 4A5E 6C76 B6DC
Ncat: bind to :::80: Permission denied. QUITTING.
```

Additional resources

- For more information, see the **udica(8)** and **podman(1)** man pages.

6.3. ADDITIONAL RESOURCES

- For more details on creating policies with **udica**, see the [udica - Generate SELinux policies for containers](#) page.

CHAPTER 7. DEPLOYING THE SAME SELINUX CONFIGURATION ON MULTIPLE SYSTEMS

This section provides two recommended ways for deploying your verified SELinux configuration on multiple systems:

- Using RHEL System Roles and Ansible
- Using **semanage** export and import commands in your scripts

7.1. INTRODUCTION TO THE SELINUX SYSTEM ROLE

RHEL System Roles is a collection of Ansible roles and modules that provide a consistent configuration interface to remotely manage multiple RHEL systems. The SELinux system role enables the following actions:

- Cleaning local policy modifications related to SELinux booleans, file contexts, ports, and logins.
- Setting SELinux policy booleans, file contexts, ports, and logins.
- Restoring file contexts on specified files or directories.

The following table provides an overview of input variables available in the SELinux system role.

Table 7.1. SELinux system role variables

Role variable	Description	CLI alternative
<code>selinux_policy</code>	Chooses a policy protecting targeted processes or Multi Level Security protection.	SELINUXTYPE in /etc/selinux/config
<code>selinux_state</code>	Switches SELinux modes. See ansible-doc selinux	setenforce and SELINUX in /etc/selinux/config .
<code>selinux_booleans</code>	Enables and disables SELinux booleans. See ansible-doc seboolean .	setsebool
<code>selinux_fcontexts</code>	Adds or removes a SELinux file context mapping. See ansible-doc sefcontext .	semanage fcontext
<code>selinux_restore_dirs</code>	Restores SELinux labels in the file-system tree.	restorecon -R
<code>selinux_ports</code>	Sets SELinux labels on ports. See ansible-doc seport .	semanage port
<code>selinux_logins</code>	Sets users to SELinux user mapping. See ansible-doc selogin .	semanage login

The `/usr/share/doc/rhel-system-roles/selinux/example-selinux-playbook.yml` example playbook installed by the **rhel-system-roles** package demonstrates how to set the targeted policy in enforcing mode. The playbook also applies several local policy modifications and restores file contexts in the `/tmp/test_dir/` directory.

Additional resources

- For a detailed reference on SELinux role variables, install the **rhel-system-roles** package, and see the **README.md** or **README.html** files in the `/usr/share/doc/rhel-system-roles/selinux/` directory.
- For more information on RHEL System Roles, see [What RHEL System Roles are and which tasks they can be used for](#)

7.2. USING THE SELINUX SYSTEM ROLE TO APPLY SELINUX SETTINGS ON MULTIPLE SYSTEMS

Follow the steps to prepare and apply an Ansible playbook with your verified SELinux settings.

Prerequisites

- Your Red Hat Ansible Engine subscription is attached to the system. See the [How do I download and install Red Hat Ansible Engine](#) article for more information.

Procedure

1. Enable the RHEL Ansible repository, for example:

```
# subscription-manager repos --enable ansible-2-for-rhel-8-x86_64-rpms
```

2. Install Ansible Engine:

```
# yum install ansible
```

3. Install RHEL system roles:

```
# yum install rhel-system-roles
```

4. Apply your playbook with an SELinux system role.

The following command applies an example playbook, which is a part of the **rhel-system-roles** package. You can use this playbook as a template:

```
# ansible-playbook -i host1,host2,host3 /usr/share/doc/rhel-system-roles/selinux/example-selinux-playbook.yml
```

Additional resources

- For more information, install the **rhel-system-roles** package, and see the `/usr/share/doc/rhel-system-roles/selinux/` and `/usr/share/ansible/roles/rhel-system-roles.selinux/` directories.

7.3. TRANSFERRING SELINUX SETTINGS TO ANOTHER SYSTEM WITH SEMANAGE

Use the following steps for transferring your custom and verified SELinux settings between RHEL 8-based systems.

Prerequisites

- The **polycoreutils-python-utils** package is installed on your system.

Procedure

1. Export your verified SELinux settings:

```
# semanage export -f ./my-selinux-settings.mod
```

2. Copy the file with the settings to the new system:

```
# scp ./my-selinux-settings.mod new-system-hostname
```

3. Log in on the new system:

```
$ ssh root@new-system-hostname
```

4. Import the settings on the new system:

```
new-system-hostname# semanage import -f ./my-selinux-settings.mod
```

Additional resources

- **semanage-export(8)** and **semanage-import(8)** man pages